

Tools for uncomfortably big data

Edward Visel
LA R Users Group Meetup
2020-09-22

Edward Visel



The short version:

- Philosophy major
- Worked in campaigns
- Data scientist at Uptake, which uses ML to identify problems in assets in the industrial internet of things
- Have worked on:
 - Anomaly detection algorithms
 - Productionalizing data science
 - Delivering DS for our Fleet/Federal team for customers with truck-like assets
- Obsessions: fried egg tacos and clean code

Uncomfortably big data

For the purposes of this talk:

- **Small data:** data you can work with locally without any concessions due to size
- **Uncomfortable data:** data you can work with locally with a little forethought
- **Big data:** data big enough to require significant tooling, probably stored in a managed database or data lake

Maxim for Uncomfortable Data:

Subset to what you need



From uncomfortable to small data

For on-disk data:

Strategy 1: Read it all in, then subset

- Simplest option for data which will fit in memory

Strategy 2: Subset on read

- Some reading functions will allow you to subset columns
 - `readr`'s `col_types` parameter
 - `data.table::fread`'s `select` and `drop` parameters

From uncomfortable to small data

Strategy 3: Read in batches, subset and aggregate

- Almost all reader function have **skip** and **nrows/n_max** params that can be used for batching
- Requires more setup
 - Have to specify column names
 - Be cautious about types so aggregation works as expected

Caution! Don't make 10,000 data frames in your global environment (😬) or try to grow a single data frame by appending during iteration (requires a lot of reallocation of memory); make a list of data frames from the start and combine at the end.

Demo 1

Subsetting strategies

[Code for demos](#)

Storing uncomfortable data

If you don't regularly **create large datasets**, you might think you can ignore how to write them...but that's wrong because:

- Better stored data can be **interacted with more efficiently**, to the point that it's often worthwhile to re-save data
- You may want to save large intermediate products that took a long time to create
- Nobody really wants to deal with a multi-gigabyte Excel file with unparsable garbage in there somewhere, or a 281 Gb CSV

Storing uncomfortable data

When to CSV:

- Your data is small enough that you can look at it all, and visual inspection is useful
- You want to use command-line tools on it
- You're giving it to someone who doesn't code

When not to CSV:

- Your data is 281 Gb
- You want to ensure types are maintained
- You care about how much space it takes up
- You often want to subset it without reading the whole thing

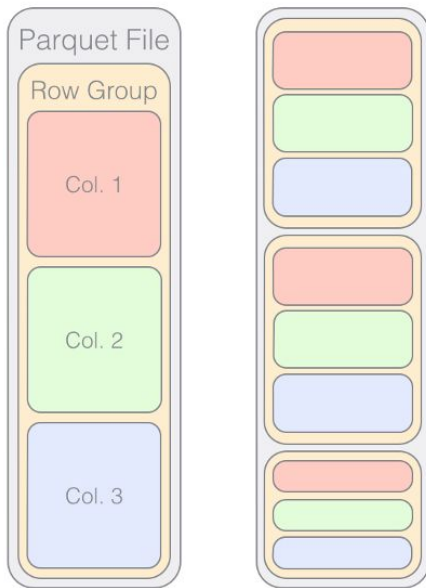
Storing uncomfortable data

Why Apache Parquet

- A serialized, compressed format like .Rds that **takes up less space** than plaintext
- Well-supported by **lots of languages** (R, Python, C++, Java, Scala, ...) and tools
- **Maintains types**, so datetimes will stay datetimes and integers will stay integers without a reader function guessing
- Optimized so individual columns can be read without scanning the whole file, so **reads are fast**

Parquet in R

in 1 minute



[Source](#)

Read: `arrow::read_parquet()`

- `col_select` parameter accepts either a vector of names to keep, or a tidy selection specification

Write: `arrow::write_parquet()`

- `version` defaults to "1.0" to maximize compatibility, but "2.0" enables some cool features
- Default `compression` is "[snappy](#)", which aims for reasonably small sizes very fast. "[gzip](#)" and others are also available, and `compression_level` can be set.
- `chunk_size` is the number of rows to put in a row group. (Think of row groups as files within the file.) When reading large files, smaller row group sizes can allow parallelism, reduce memory consumption, and when [well optimized](#) allow groups not to be read if irrelevant.

Storing uncomfortable data

Save one big file or lots of small files?

- One big file allows more familiar idioms—if the machine can handle it
- **More smaller files** are less likely to blow out the machine's resources
 - Partitioning makes **parallelization** easy
 - Tooling can make handling a directory of files as simple as a single file

Before splitting a dataset, consider your **partitioning scheme**

- How will the dataset be read? What **natural divisions** when used for partitioning will allow only the desired data to be read?

Demo 2

Storing data

Tools for working with on-disk data

When working with data on-disk that may be bigger than available memory, batching like in **Strategy 3** still works...but still takes time to write and optimize.

Tooling can help! Good tools let you *subset to what you need **faster***. Today's (non-exhaustive!) subjects:

- [Arrow Datasets](#) are a quick, **lightweight** way to handle on-disk, partitioned data
- [Apache Drill](#) is a robust tool that allows on-disk data to be **queried with SQL** like a database

Tools for working with on-disk data

[Apache Arrow](#) is a big project focused on standardizing how data is represented in-memory.

Arrow Datasets are a relatively new tool in R and Python to work with directories of on-disk data.

- Originally focused around **Parquet**, but now supports other formats including **CSV**
- Supports a very limited subset of **dplyr syntax**
- Great for *subsetting to what you need*; not great for calculating summary stats or sophisticated operations

Tools for working with on-disk data

To create an Arrow Dataset, use `arrow::open_dataset()`

- Pass `sources` a path to a directory of data files
- Use `format` to specify filetype. Default is `"parquet"`; other options include `"csv"`, `"tsv"`, and `"feather"`
- `partitioning` can be used to capture directory structure within the `source` path as variables
 - E.g. with a relative path of `2020/09/22/data.parquet` and `partitioning = c("year", "month", "day")`, three variables of those names will be appended to the dataset

Tools for working with on-disk data

Arrow Dataset dplyr syntax

- Accepts `select()`, `rename()`, `filter()` with very simple conditions, and `group_by()`
- Delayed evaluation, so use `collect()` to evaluate and get data, like working with a database from dplyr
- Does not yet accept `mutate()` or `summarise()` calls until after collection (at which point you are working with a normal, in-memory data frame)

👉 Great for *subsetting to what you need*; not great for your whole workflow

Demo 3

Arrow Datasets

Apache Drill

tl;dr

- Drill lets you write SQL against files without standing up a database
 - The files can be anywhere, including your hard drive
 - Lots of file formats are supported (CSV, TSV, Parquet, JSON, etc.)
 - Globbing works
 - e.g. `ls /tmp/*.csv`

What is it?

Schema-free **SQL Query Engine** for Hadoop, NoSQL and Cloud Storage

Drill supports a variety of NoSQL databases and file systems, including HBase, MongoDB, MapR-DB, HDFS, MapR-FS, Amazon S3, Azure Blob Storage, Google Cloud Storage, Swift, NAS and **local files**. A single query can join data from multiple datastores.



SQL on local files!

Installing Drill and accessories

Install Drill via a package manager on your OS. For macOS with homebrew,

```
$ brew install apache-drill
```

- Drill requires Java 1.8.
- To keep Drill available as a service, install [zookeeper](#) similarly and start it ([instructions](#)).

To use Drill from R, install [sergeant](#) (uses ODBC interface) or [sergeant.caffeinated](#) (uses JDBC; requires working rJava) as usual with `install.packages()`.

Starting Drill

To run Drill locally in embedded mode (not as a service), after installation start it from the command line with

```
$ drill-embedded
```

and you'll be greeted by a **Drill shell** where you can run queries and commands.

Once Drill is running, navigate to <http://localhost:8047/> (adjust port if you've changed it) to see a website-based interface where you can query, configure, and administer.

Querying a filesystem with Drill

Storage plugins tell Drill where to look

- The local filesystem is available as `dfs`
- Can configure to change file handling behavior. Other storage plugins can be added, e.g `s3`

Workspaces are path shortcuts

- Defaults for `dfs` include `dfs.root (/)` and `dfs.tmp (/tmp)`. Add your own!

Pass `FROM` clause of a query

1. Storage plugin, e.g. `dfs`
2. Workspace joined by `.`, e.g. `tmp`
3. Path, wrapped in ``...``

...e.g. `dfs.tmp.`flights/*``

Handy features

Use **SHOW FILES FROM ...** to navigate

Set a default workspace with **USE**

- E.g. after **USE dfs.tmp;** can use **FROM `flights`** instead of **FROM dfs.tmp.`flights`**

Implicit columns offer access to path and filename

- Select **FQN** (fully qualified name), **FILEPATH**, **FILENAME**, or **SUFFIX**, and they'll magically appear as columns
- Subdirectories of the directory query will be appended as **dir0**, **dir1**, ...

Demo 4

Apache Drill

...ok cool but I thought this was a talk
about tools I could use in R



Connect to Drill from R with sergeant

sergeant is an R package by Bob Rudis which defines a DBI interface and dbplyr backend for Drill

- Drill in a docker container with `drill_up()`
- Has both Drill **ODBC** and **REST** interfaces
- To use **JDBC**, use `sergeant.caffeinated`, which works the same way
- To use dbplyr, pass `tbl()` a DBI connection created with `dbConnect(Drill())` or `src_drill()` and a **FROM** clause argument
- Lots of nice helpers for SQL translation! See `?drill_custom_functions` and `sql_translate_env(src_drill())$con`

dbplyr

basics

dbplyr: A database backend for dplyr

- Most dplyr verbs (e.g. `select`, `filter`, `mutate`, `summarise`) work as usual
- Operations in mutating verbs may not.
 - Recognized functions will be translated to SQL
 - Unrecognized ones will be passed through to Drill untranslated, allowing the use of SQL functions
- Queries are lazy and only return the first 10 rows. To return everything, run `collect()`
- To see what query will be run, use `show_query()`

Demo 5

Drill in R via sergeant

References and reading

- [Querying across files with Apache Drill](#)
- Drill docs
 - [Drill in 10 Minutes](#)
 - [Drill SQL reference](#)
- [sergeant](#)
 - [sergeant.caffeinated](#)
 - [Using Apache Drill with R](#) ← start here
- [dbplyr](#)
 - [Getting started with dbplyr](#)
 - [Writing SQL with dbplyr](#)
 - [Function translation](#)
- [Arrow](#) and [parquet](#)
 - [arrow R package](#)
 - [Datasets vignette](#)

